

# CSEN 166 Artificial Intelligence

## Lab Assignment #5 Reinforcement Learning

This is a **GROUP** assignment. You will download **rl.zip** and modify the functions in “class **QLearningAgent(ReinforcementAgent):**” of **qlearningAgents.py**

**Task A:** Implement a Q-learning agent to walk in the grid world we saw in class. The agent learns by trial and error from interactions with the environment through its **update(state, action, nextState, reward)** method. A stub of a Q-learner is specified in **QLearningAgent** in **qlearningAgents.py**, and you can select it with the option **'-a q'**.

You will implement the **update**, **computeValueFromQValues**, **getQValue**, and **computeActionFromQValues** functions in **qlearningAgents.py**.


*Note:* For **computeActionFromQValues**, you should break ties randomly for better behavior. The **random.choice()** function will help. In a particular state, actions that your agent *hasn't* seen before still have a Q-value, specifically a Q-value of zero.

*Important:* Make sure that in your **computeValueFromQValues** and **computeActionFromQValues** functions, you only access Q values by calling **getQValue**.

**Task A – Experiment 1** With the Q-learning update in place, you can watch your Q-learner learn under manual control, using the keyboard:

**python gridworld.py -a q -k 4 -m --noise 0.0**

Note: **-k** will control the number of episodes your agent gets to learn, **-m** means you manually control Pacman's actions, and **--noise 0.0** means there is no randomness (it's a deterministic game). Please manually steer Pacman north and then east along the optimal path for **four episodes**, and you should see the following updated Q-values:



Hint:

- In `computeValueFromQValues` and `computeActionFromQValues`, the available actions shall be obtained by: `self.getLegalActions(state)`
- Watch how the agent learns about the state it was just in, not the one it moves to.

**Task A – Experiment 2** Run the command:

```
python autograder.py -q q3
```

Your code is expected to pass all four test\_cases of the above command.

**Task B:** Keep the modified code in Task A. Now, you will continue to modify the “`def getAction(self, state):`” function in the same `qlearningAgents.py` file. This function computes the action to take in the current state. There is a parameter `self.epsilon`, which is a probability. With probability `self.epsilon`, the agent should take a random action of all legal actions available in that state. With probability `1 - self.epsilon`, the agent will take the best policy action obtained by the `computeActionFromQValues` function.

**Task B – Experiment 1** Run the following command, and observe how the Q-learner updates the Q-state values for 100 episodes.

```
python gridworld.py -a q -k 100
```

**Task B – Experiment 2** Run the following command and observe how the crawler bot learns to move to the right. You can decrease the “Step Delay” to 0.0125 by pressing the top-left “-” button, so that you will see the displayed results sooner.

```
python crawler.py
```

**Task B – Experiment 3** Run the following autograder, and your code should pass all four test\_cases.

```
python autograder.py -q q4
```

### Submission:

1. Submit a pdf file to Camino (the format of the file is similar to Lab4\_submission\_sample.pdf. Also include a screen shot/picture of the results you get for each experiment. For instance, for the crawler bot experiment, include a picture of the crawler bot when it's in the middle of the scene.).
2. Submit all source code needed (with qlearningAgents.py modified by you) to generate the results of **all Experiments in Tasks A and B** as a .zip file to Camino. We will test run your submitted code, so make sure it works.

**Demonstration:** you will demonstrate all experiments in tasks A and B to the TA, and explain how you implemented the following functions in your qlearningAgents.py code:

- getQValue
- computeValueFromQValues
- computeActionFromQValues
- update
- getAction

These functions are essential components of your algorithm, and you will also explain them in your lab report.

**Grading:** All demonstrations must be done before the lab report is due. The grade depends on both the correctness of the demonstration and submitted report and code.